Name: Pradeep Medagiri

Homework 10: Lambda

You are building an information system for a startup financial company in Northern Virginia. It needs a database to store companies' financial information, such as their stock price, industry they are in, and some other information (some companies have more information than others). This database is intended for the customers so that they can **query** information quickly using company's **ticker**. You decide to use the NoSQL and fast AWS **DynamoDB** for the database.

As part of their work process, a staff will research the financial information of a company, create a json file to record the data, and upload the json file in S3. One company per json file.

The company wants to implement Robotic Process Automation (RPA). In other words, they want to automate repetitive activities with software. So instead of populating the database manually, they want the system to automatically read financial data every time a json file is uploaded to the S3 bucket, and to populate the data into the database accordingly.

You need to create a lambda function (similar to the demo) to automatically read uploaded json file and populate the dynamodb database.

You have data of 4 companies and you need to create a json file for each of them for testing purposes (all numbers are integers to simplify things).

Ticker: AAPL
Stock_price: 389
Industry: Computer
Name: Apple
Volume: 20279276

Ticker: IBM
Stock_price: 129
Name: International Business Machines

Ticker: HPE
Stock_price: 10
Industry: Computer
Name: Hewlett Packard Enterprise

Ticker: TSLA
Stock_price: 1592
Industry: Automobile
Name: Tesla
Volume: 14161080
State: California
City: San Francisco

Your tasks:

1. Create a dynamodb table – name it yourlastname-hw10-dynamodb.
2. Create an S3 bucket with private access to store the json files.
3. Create a json file for each of the companies listed above. Name the each of the json file using the ticker. For example, the json file for the apple is AAPL.json
4. Create a lambda function that will read json file uploaded to the S3 and automatically populate the dynamodb table.
5. Upload all four json files. All four companies must be in the dynamodb database.

What to submit:

1. The screenshot of the IAM Role you created that shows all its policies. (16%)

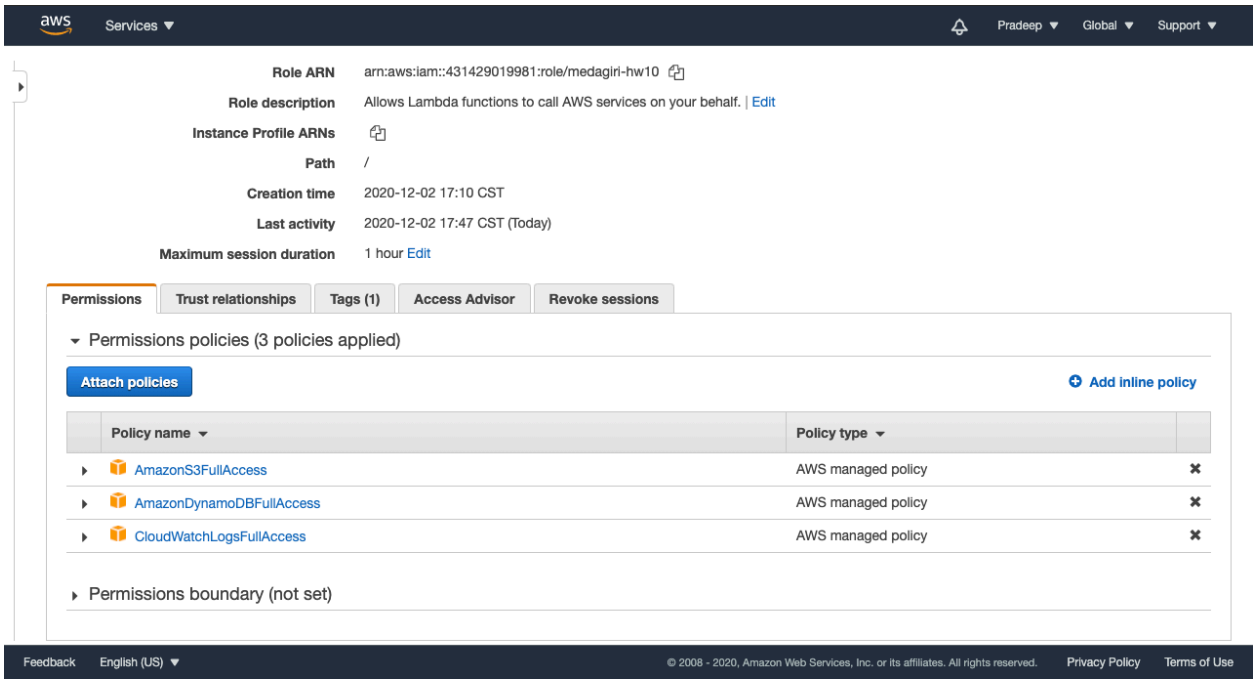

2. The screenshot of the final working lambda function. (20%)

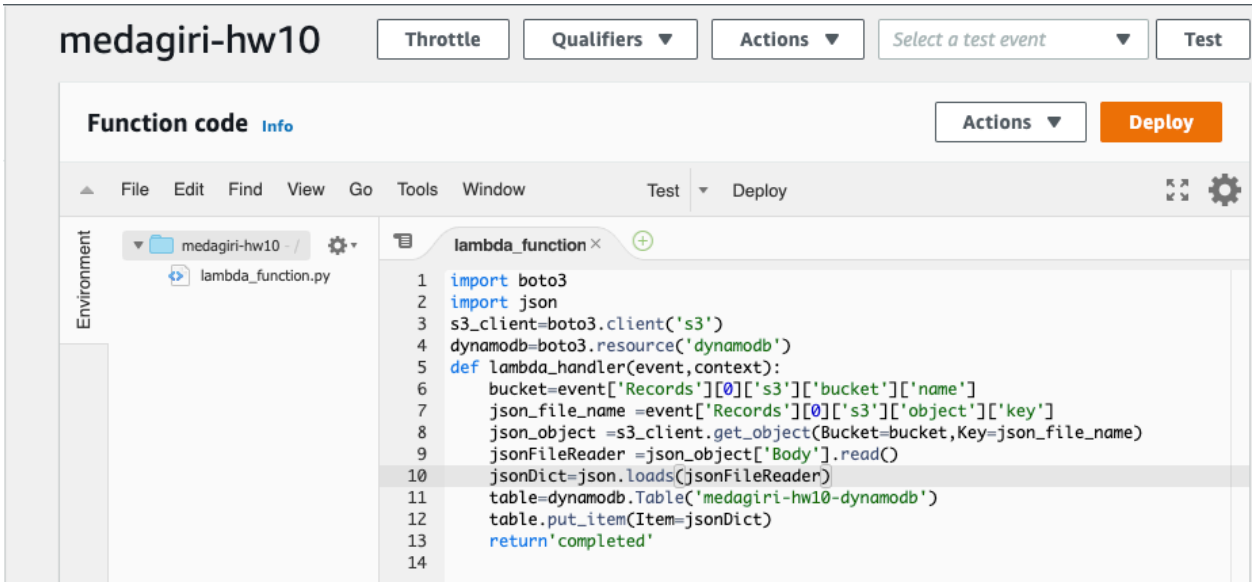

3. The screenshot of Tesla json file. (16%)

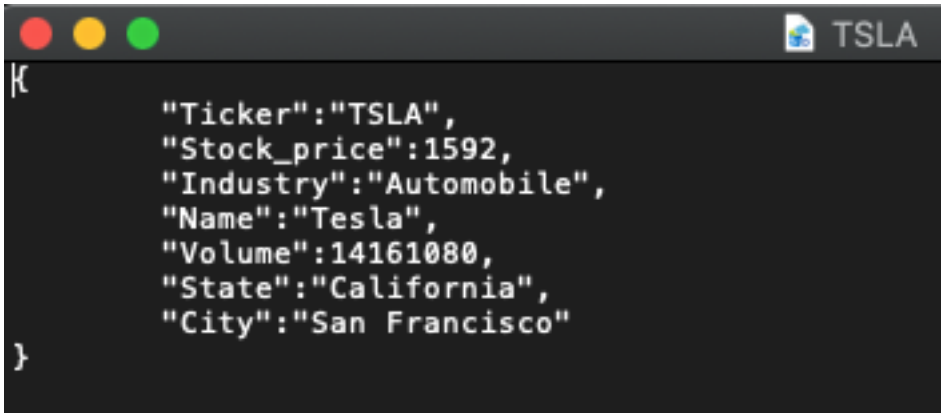

4. The screenshot of the S3 bucket with all the json files uploaded. (16%)

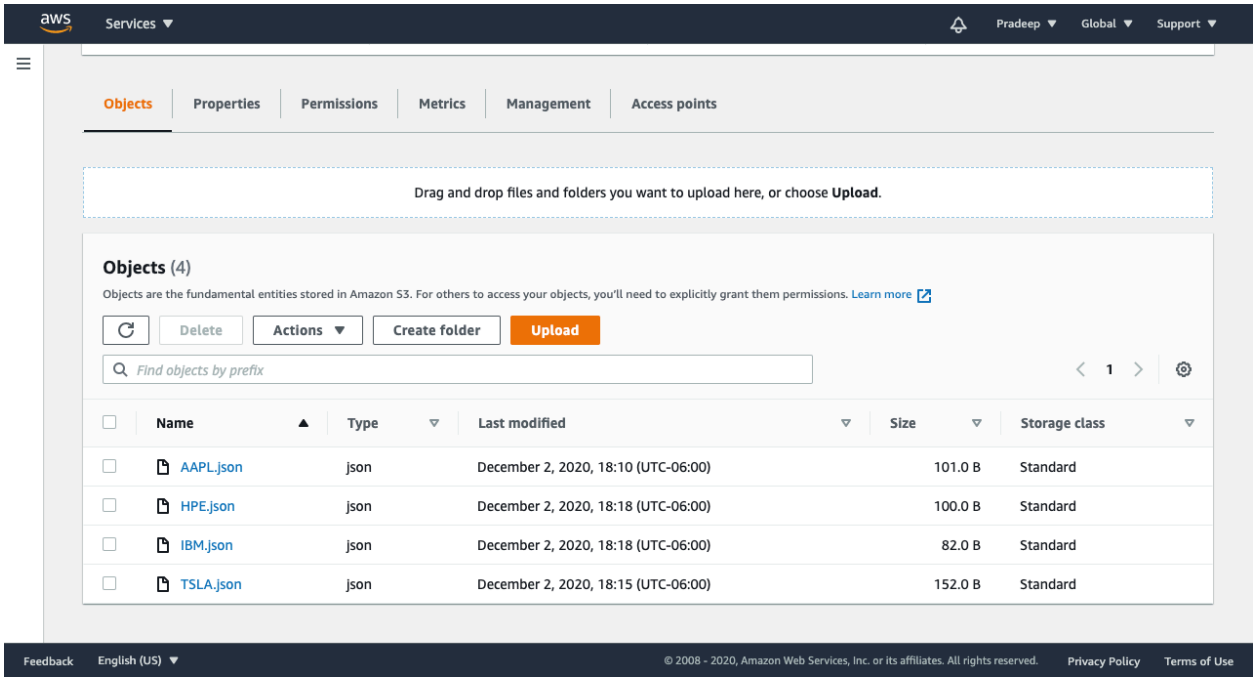

5. The screenshot CloudWatch log groups for this lambda function that shows no errors (after all 4 json files have been uploaded). (16%)

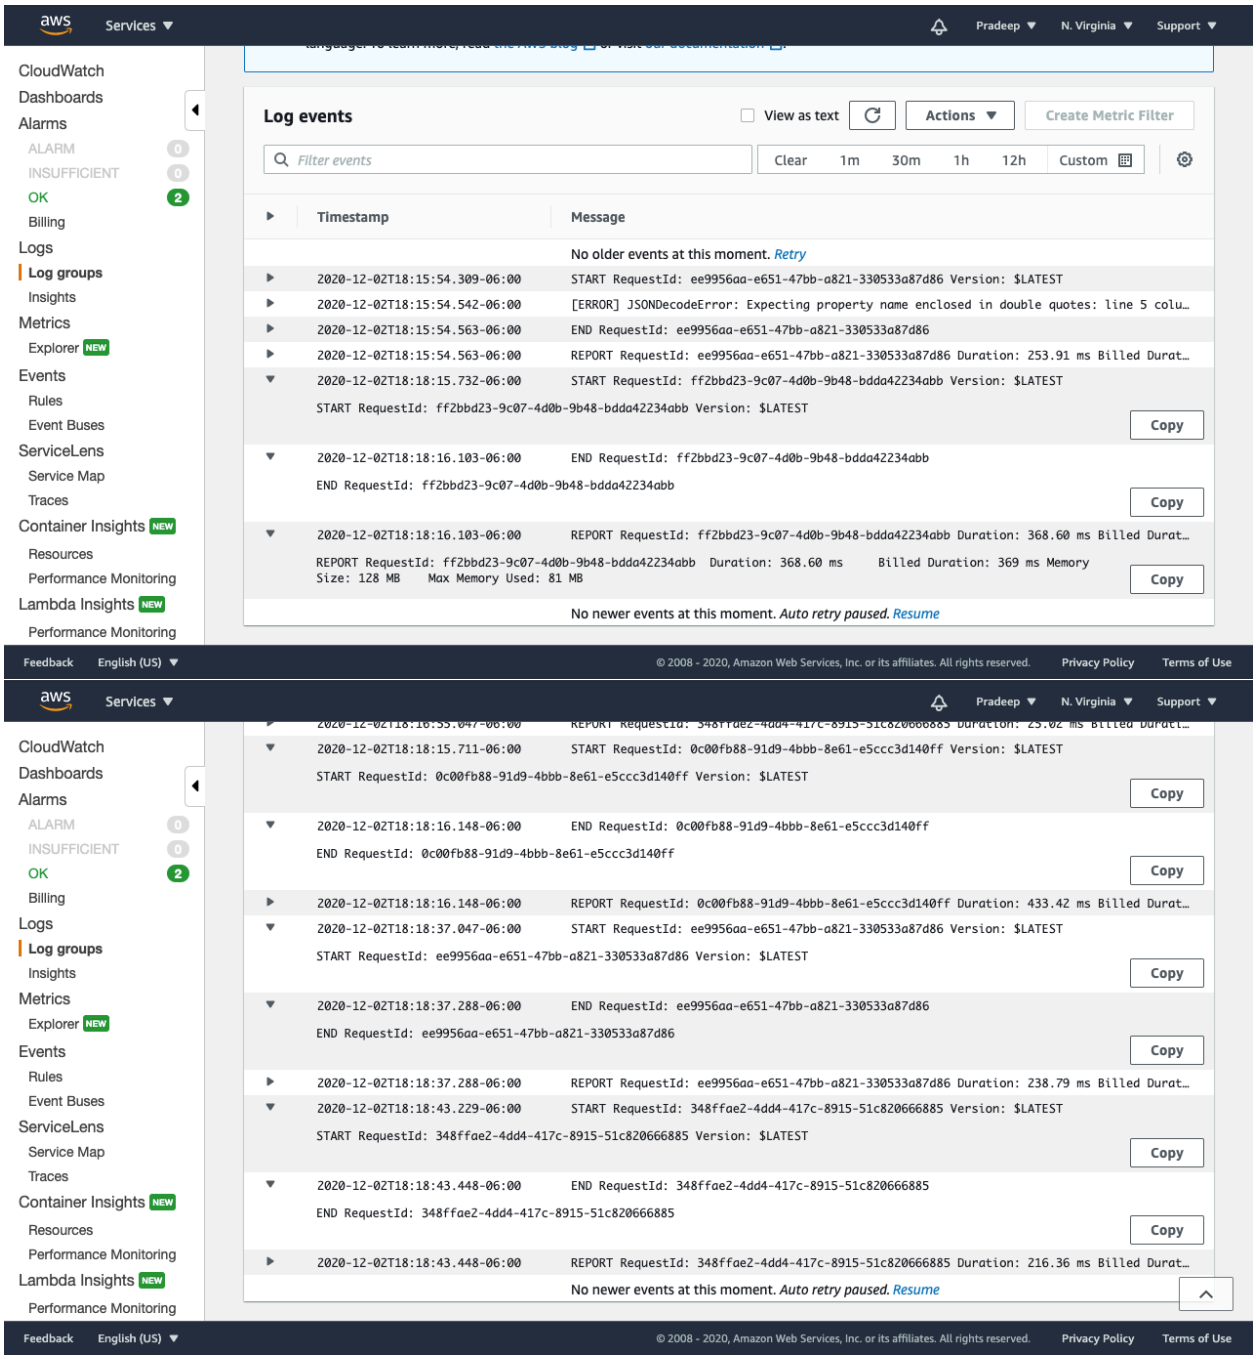

6. The screenshot of the items in the **dynamodb table (after all 4 json files have been uploaded)**. (16%)