Name: Pradeep Medagiri

You are assigned to create a NoSQL key-value database for a product catalog using AWS DynamoDB for a new store located in Northern Virginia. As is normally done in development stage, you start with a handful of items to test if the database works. Here are the parameters:

Table name: yourlastnameProductCatalog

Primary Key:

- Partition Key: ID (attribute type: number)

Global Secondary Index

- ProductCategory (attribute type: string)

Use Cloud9 or whatever IDE you prefer and create python scripts to:

- Create the dynamoDB table
- Define the Primary Key and Global Secondary Index
- To insert the 6 items below
- To query item ID number 103
- To query all items in the Bicycle category

What to submit:

1. The screenshot of the DynamoDB table creation python script and show that it runs successfully (show the file editor and the terminal). (5)
2. The screenshot of the Global Secondary Index definition python script and show that it runs successfully (show the file editor and the terminal). (5)
3. The screenshot of the item insertion python script and show that it runs successfully (show the file editor and the terminal). (5)
4. The screenshot of the query item ID number 103 python script and show the result (file editor and terminal). (5)
5. The screenshot of the query all items in the Bicycle category python script and show the result (file editor and terminal. (5)

Example of screenshot:

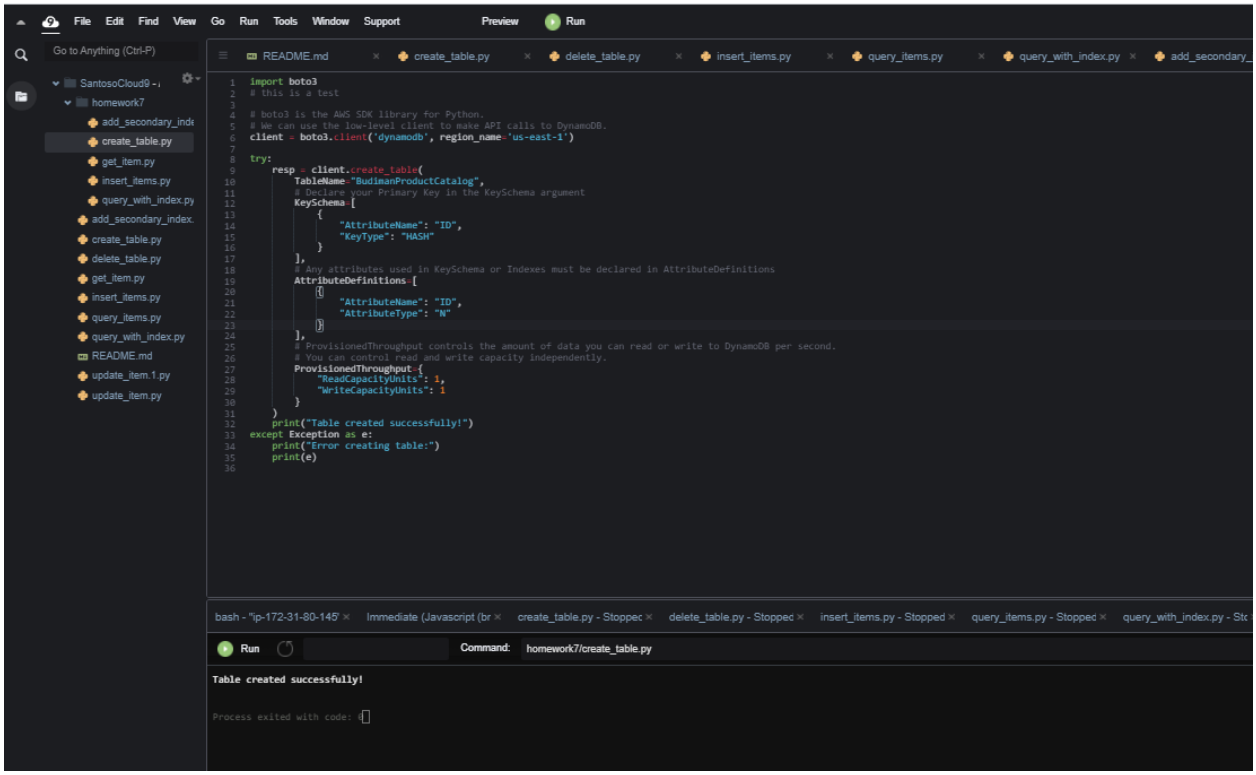

Note: **if there is no " ", the value is not a string**

1st Item

"Id": 101
"Title": "Book 101 Title"
"ISBN":"111-1111111111"
"Authors": "Author1"
"Price": 2
"Dimensions": "8.5 x 11.0 x 0.5"
"PageCount": 500
"InPublication": true
"ProductCategory": "Book"

2nd item

"Id": 102
"Title": "Book 102 Title"
"ISBN": "222-2222222222"
"Authors": "Author2
"Price": 20
"Dimensions": "8.5 x 11.0 x 0.8"
"PageCount": 600
"InPublication": true
"ProductCategory": "Book"

3rd item

"Id": 103
"Title": "Book 103 Title"
"ISBN": "333-3333333333"
"Authors": "Author1"
"Price": 2000
"Dimensions": "8.5 x 11.0 x 1.5"
"PageCount": 600
"InPublication": false
"ProductCategory": "Book"

4th item

"Id": 201
"Title": "18-Bike-201"
"Description": "201 Description"
"BicycleType": "Road"
"Brand": "Mountain A"
"Price": 100
"Color": ["Red",”Black”]
"ProductCategory": "Bicycle"

5th item

"Id": 202
“Title": "21-Bike-202"
"Description": "202 Description"
"BicycleType": "Road"
"Brand": "Brand-Company A"
"Price": 200
"Color": ["Green", "Black"]
"ProductCategory": "Bicycle"

6th item

"Id": 203
"Title": "19-Bike-203"
"Description": "203 Description"
"BicycleType": "Road"
"Brand": "Brand-Company B"
"Price": 300
"Color": ["Red","Green","Black"]
"ProductCategory": "Bicycle"

Answers:

What to submit:

1. The screenshot of the DynamoDB table creation python script and show that it runs successfully (show the file editor and the terminal). (5)

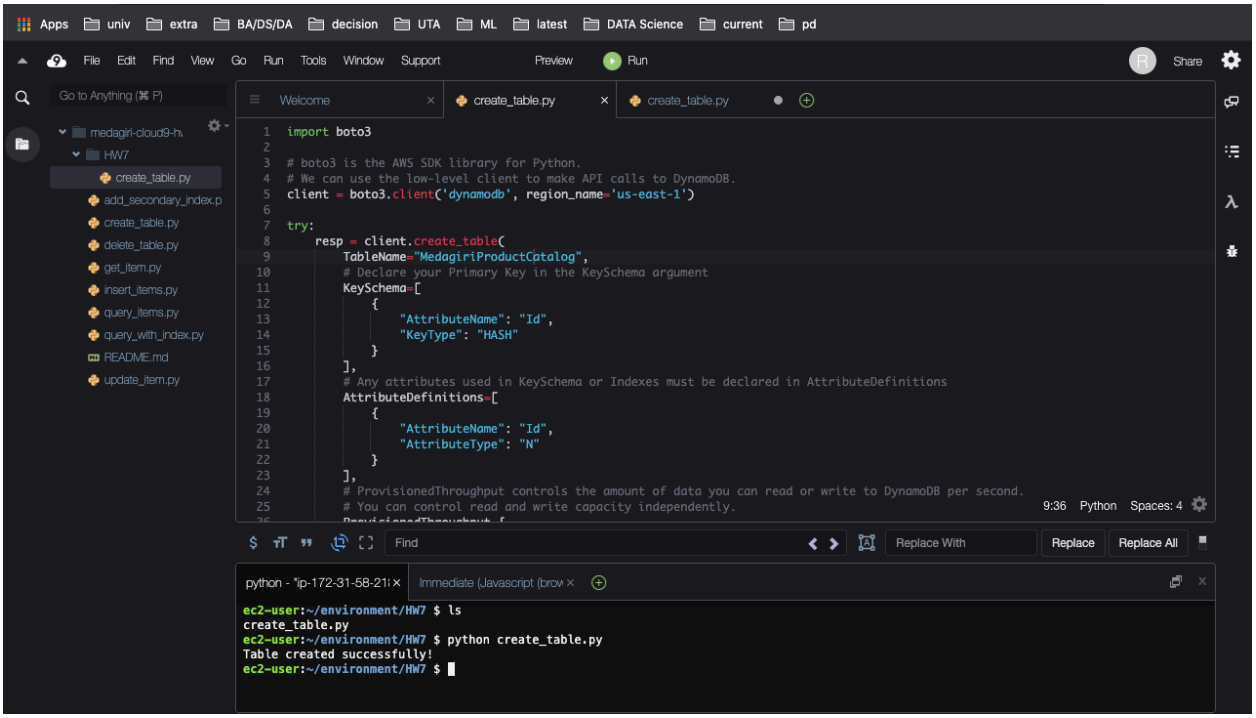

2. The screenshot of the Global Secondary Index definition python script and show that it runs successfully (show the file editor and the terminal). (5)

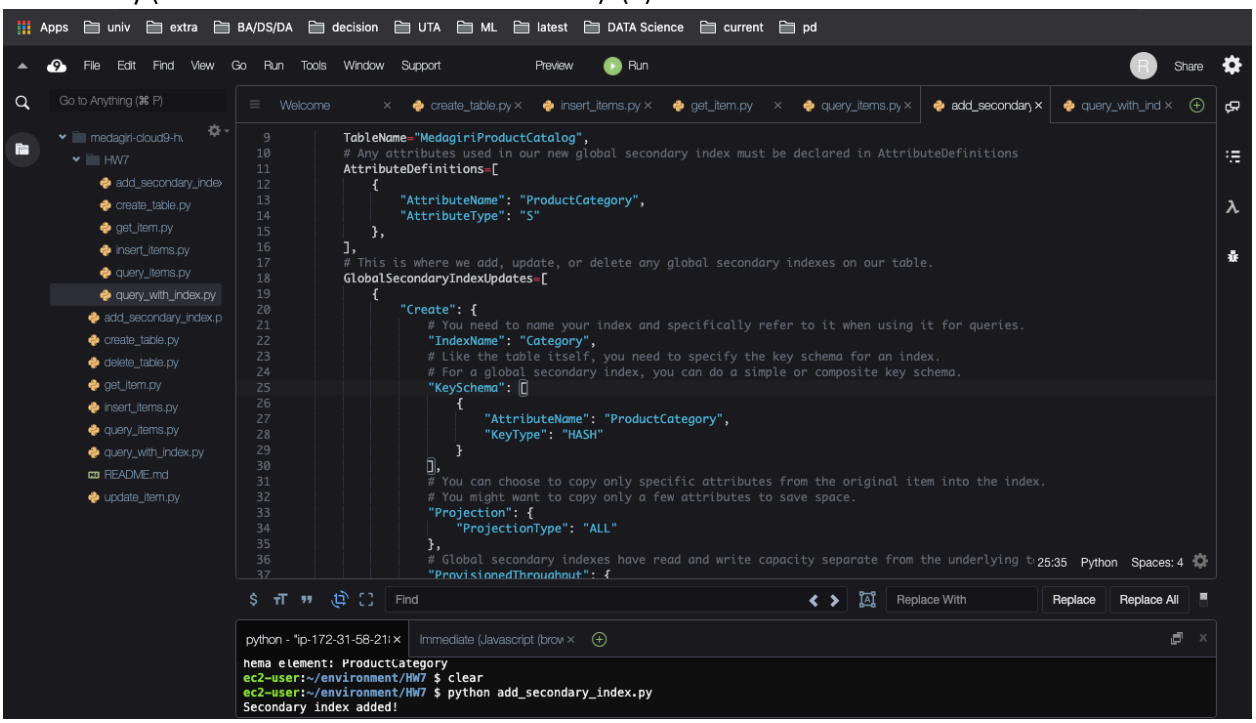

3. The screenshot of the item insertion python script and show that it runs successfully (show the file editor and the terminal). (5)

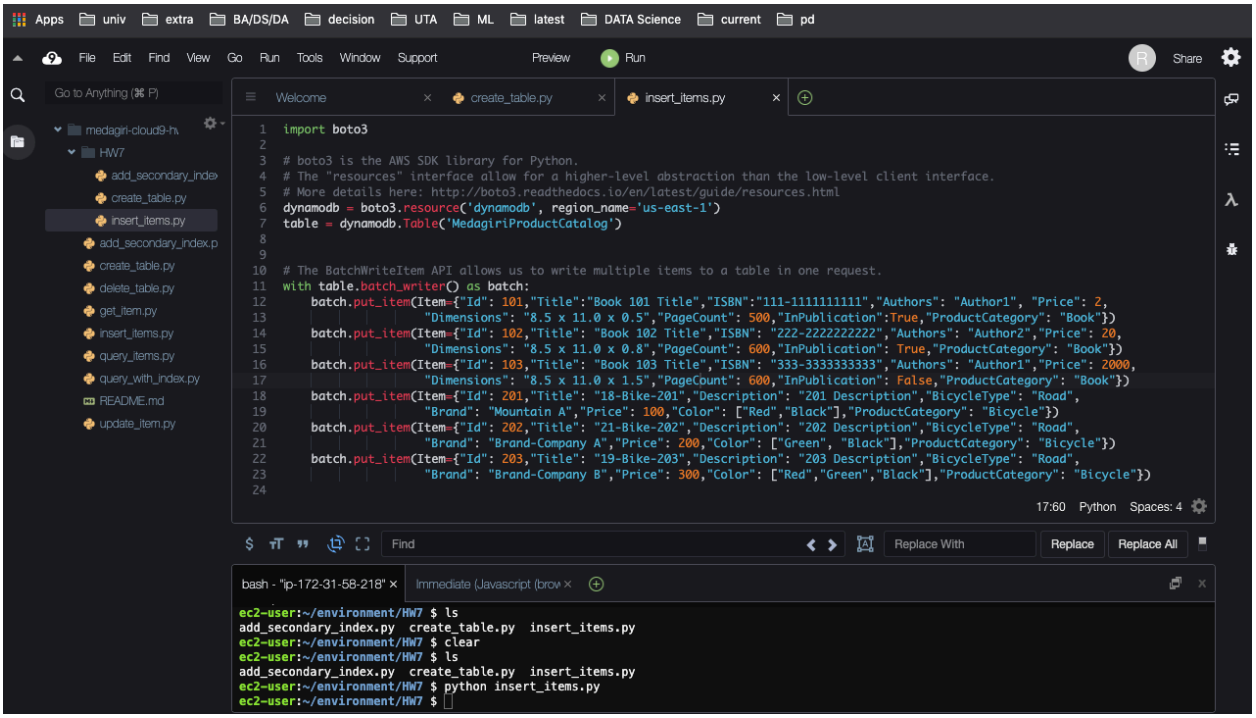

4. The screenshot of the query item ID number 103 python script and show the result (file editor and terminal). (5)

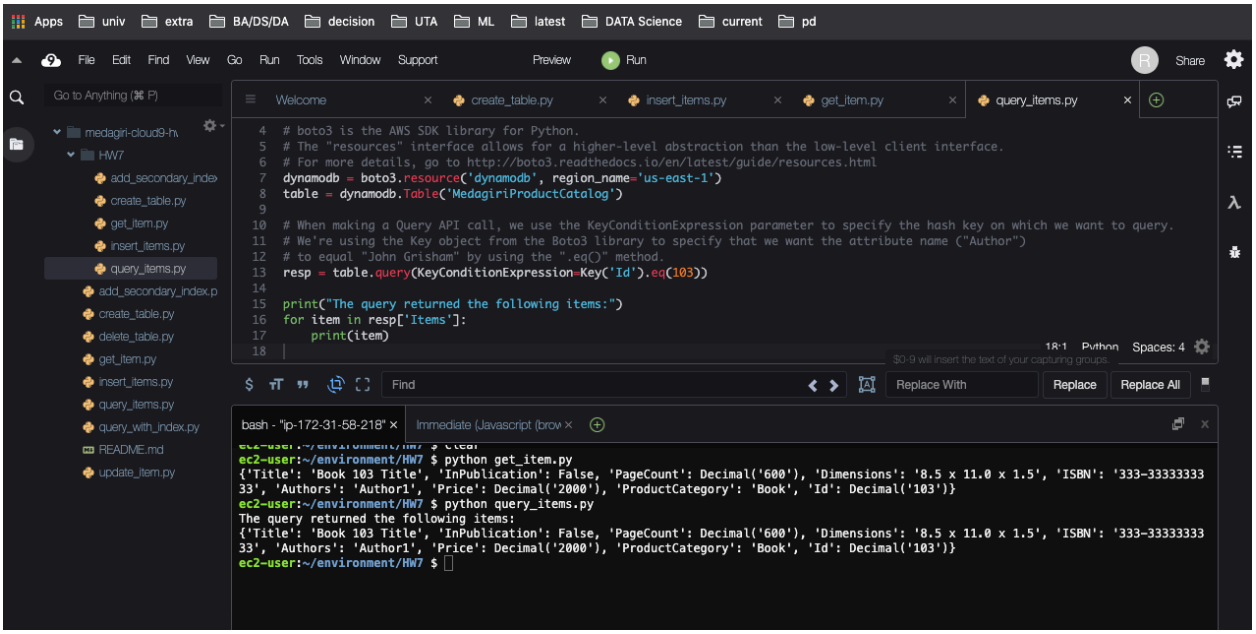

5. The screenshot of the query all items in the Bicycle category python script and show the result (file editor and terminal. (5)